



# Multiprocessor and Multicore Scheduling

# Coarse and Very Coarse Grained Parallelism

---

- Synchronization among processes, but at a very gross level
- Good for concurrent processes running on a multiprogrammed uniprocessor
  - can be supported on a multiprocessor with little or no change to user software



# Assignment of Processes to Processors

Assuming all processors are equal, it is simplest to treat processors as a pooled resource and assign processes to processors on demand

static or dynamic needs to be determined

If a process is permanently assigned to one processor from activation until its completion, then a dedicated short-term queue is maintained for each processor

advantage is that there may be less overhead in the scheduling function

allows group or gang scheduling

- A disadvantage of static assignment is that one processor can be idle, with an empty queue, while another processor has a backlog
  - to prevent this situation, a common queue can be used
  - another option is dynamic load balancing

# Process Scheduling

---

- Usually processes are not dedicated to processors
- A single queue is used for all processors
  - if some sort of priority scheme is used, there are multiple queues based on priority
- System is viewed as being a multi-server queuing architecture



# Load Sharing

- Simplest approach and carries over most directly from a uniprocessor environment

## Advantages:

- load is distributed evenly across the processors
- no centralized scheduler required
- the global queue can be organized and accessed using any of the schemes discussed in Chapter 9

- Versions of load sharing:
  - first-come-first-served
  - smallest number of threads first
  - preemptive smallest number of threads first



# Dedicated Processor Assignment

---

- When an application is scheduled, each of its threads is assigned to a processor that remains dedicated to that thread until the application runs to completion
- If a thread of an application is blocked waiting for I/O or for synchronization with another thread, then that thread's processor remains idle
  - there is no multiprogramming of processors
- Defense of this strategy:
  - in a highly parallel system, with tens or hundreds of processors, processor utilization is no longer so important as a metric for effectiveness or performance
  - the total avoidance of process switching during the lifetime of a program should result in a substantial speedup of that program

# Cache Sharing

---

## Cooperative resource sharing

- Multiple threads access the same set of main memory locations
- Examples:
  - applications that are multithreaded
  - producer-consumer thread interaction

## Resource contention

- Threads, if operating on adjacent cores, compete for cache memory locations
- If more of the cache is dynamically allocated to one thread, the competing thread necessarily has less cache space available and thus suffers performance degradation
- Objective of contention-aware scheduling is to allocate threads to cores to maximize the effectiveness of the shared cache memory and minimize the need for off-chip memory accesses

# Periodic and Aperiodic Tasks

---

## ■ Periodic tasks

- requirement may be stated as:
  - once per period  $T$
  - exactly  $T$  units apart

## ■ Aperiodic tasks

- has a deadline by which it must finish or start
- may have a constraint on both start and finish time



# User Control



- Generally much broader in a real-time operating system than in ordinary operating systems
- It is essential to allow the user fine-grained control over task priority
- User should be able to distinguish between hard and soft tasks and to specify relative priorities within each class
- May allow user to specify such characteristics as:

paging or  
process  
swapping


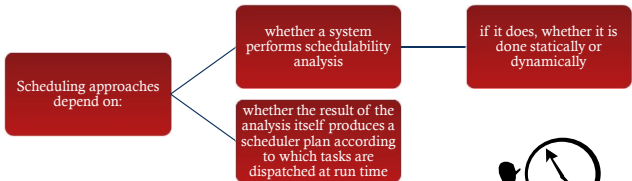
what processes  
must always be  
resident in main  
memory

what disk  
transfer  
algorithms are  
to be used

what rights the  
processes in  
various priority  
bands have

# Real-Time Scheduling

---



## Execution Profile of Two Periodic Tasks

Process	Arrival Time	Execution Time	Ending Deadline
A(1)	0	10	20
A(2)	20	10	40
A(3)	40	10	60
A(4)	60	10	80
A(5)	80	10	100
•	•	•	•
•	•	•	•
•	•	•	•
B(1)	0	25	50
B(2)	50	25	100
•	•	•	•
•	•	•	•
•	•	•	•

Processor Utilization: are both periodic tasks schedulable in the long run?

High

Priority


Low

Highest rate and  
highest priority task

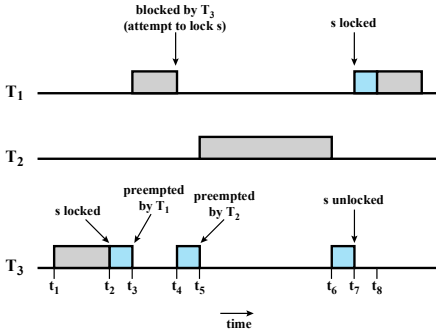
Lowest rate and  
lowest priority task

Rate (Hz)

A Task Set with RMS



# Unbounded Priority Inversion



(a) Unbounded priority inversion